

Murach Chapter 4 Part 4

How to Retrieve Data From Two or More Tables

Week 4, Lecture 7

Knowledge Points in this lecture

- OUTER JOIN of 3 tables
- Mix of OUTER JOIN and INNER JOIN
- Explicit/Implicit syntax for JOINS
 - CROSS JOIN, INNER JOIN, OUTER JOIN
- JOIN USING
- NATURAL JOIN
- Set Operators in SQL
 - UNION, INTERSECT, MINUS
- In-Class Practice – Outer Join

Query that uses 2 left outer joins of 3 tables

```
SELECT department_name, last_name, project_number AS proj_no
FROM departments d
  LEFT JOIN employees e
    ON d.department_number = e.department_number
  LEFT JOIN projects p
    ON e.employee_id = p.employee_id
ORDER BY department_name, last_name, project_number
```

| | DEPARTMENT_NAME | LAST_NAME | PROJ_NO |
|---|-----------------|-----------|---------|
| 1 | Accounting | Hernandez | P1011 |
| 2 | Maintenance | Hardy | (null) |
| 3 | Operations | (null) | (null) |
| 4 | Payroll | Aaronsen | P1012 |
| 5 | Payroll | Simonian | P1012 |
| 6 | Payroll | Smith | P1012 |
| 7 | Personnel | Jones | (null) |
| 8 | Personnel | O'Leary | P1011 |

Query that uses 2 right outer joins of 3 tables

```
SELECT department_name, last_name, project_number AS proj_no
FROM departments d
      RIGHT JOIN employees e
        ON d.department_number = e.department_number
      RIGHT JOIN projects p
        ON e.employee_id = p.employee_id
ORDER BY department_name, last_name, project_number
```

| | DEPARTMENT_NAME | LAST_NAME | PROJ_NO |
|---|-----------------|-----------|---------|
| 1 | Accounting | Hernandez | P1011 |
| 2 | Payroll | Aaronsen | P1012 |
| 3 | Payroll | Simonian | P1012 |
| 4 | Payroll | Smith | P1012 |
| 5 | Personnel | O'Leary | P1011 |
| 6 | (null) | Locario | P1013 |
| 7 | (null) | Watson | P1013 |
| 8 | (null) | (null) | P1014 |

Query that uses 2 full outer joins of 3 tables

```
SELECT department_name, last_name, project_number AS proj_no  
FROM departments dpt
```

```
    FULL JOIN employees emp
```

```
        ON dpt.department_number = emp.department_number
```

```
    FULL JOIN projects prj
```

```
        ON emp.employee_id = prj.employee_id
```

```
ORDER BY department_name
```

| | DEPARTMENT_NAME | LAST_NAME | PROJ_NO |
|----|-----------------|-----------|---------|
| 1 | Accounting | Hernandez | P1011 |
| 2 | Maintenance | Hardy | (null) |
| 3 | Operations | (null) | (null) |
| 4 | Payroll | Simonian | P1012 |
| 5 | Payroll | Aaronsen | P1012 |
| 6 | Payroll | Smith | P1012 |
| 7 | Personnel | Jones | (null) |
| 8 | Personnel | O'Leary | P1011 |
| 9 | (null) | Locario | P1013 |
| 10 | (null) | (null) | P1014 |
| 11 | (null) | Watson | P1013 |

A SELECT statement with an outer and inner join

```
SELECT department_name AS dept_name,  
       last_name, project_number  
FROM departments dpt  
       JOIN employees emp  
       ON dpt.department_number = emp.department_number  
       LEFT JOIN projects prj  
       ON emp.employee_id = prj.employee_id  
ORDER BY department_name
```

The result set

| | DEPT_NAME | LAST_NAME | PROJECT_NUMBER |
|---|-------------|-----------|----------------|
| 1 | Accounting | Hernandez | P1011 |
| 2 | Maintenance | Hardy | (null) |
| 3 | Payroll | Simonian | P1012 |
| 4 | Payroll | Smith | P1012 |
| 5 | Payroll | Aaronsen | P1012 |
| 6 | Personnel | Jones | (null) |
| 7 | Personnel | O'Leary | P1011 |

The explicit syntax for an outer join

```
SELECT select_list
FROM table_1
    {LEFT|RIGHT|FULL} [OUTER] JOIN table_2
    ON join_condition_1
[ {LEFT|RIGHT|FULL} [OUTER] JOIN table_3
  ON join_condition_2] ...
```

- [A]: A is optional
- {A|B}: A or B

What outer joins do

| Join | Keeps unmatched rows (Extra information) from |
|-------|---|
| Left | The left table |
| Right | The right table |
| Full | Both tables |

How to code a cross join with the implicit syntax

The implicit syntax for a cross join

```
SELECT select_list  
FROM table_1, table_2
```

Implicit cross join:

- Omit keyword JOIN
- No WHERE condition

A cross join that uses the implicit syntax

```
SELECT departments.department_number, department_name,  
       employee_id, last_name  
FROM departments, employees  
ORDER BY departments.department_number
```

The result set

| | DEPARTMENT_NUMBER | DEPARTMENT_NAME | EMPLOYEE_ID | LAST_NAME |
|---|-------------------|-----------------|-------------|-----------|
| 1 | 1 | Accounting | 4 | Hernandez |
| 2 | 1 | Accounting | 3 | Simonian |
| 3 | 1 | Accounting | 9 | Locario |
| 4 | 1 | Accounting | 8 | O'Leary |
| 5 | 1 | Accounting | 7 | Hardy |
| 6 | 1 | Accounting | 6 | Watson |
| 7 | 1 | Accounting | 5 | Aaronsen |

Explicit CROSS JOIN Equivalent to Previous Slide

- Use keyword: CROSS JOIN
- No WHERE condition

```
SELECT departments.department_number,  
       department_name, employee_id, last_name  
FROM departments CROSS JOIN employees  
ORDER BY departments.department_number;
```

Implicit syntax that inner-joins two tables

```
SELECT invoice_number, vendor_name
FROM vendors v, invoices i
WHERE v.vendor_id = i.vendor_id
ORDER BY invoice_number
```

The result set

| | INVOICE_NUMBER | VENDOR_NAME |
|---|----------------|-----------------------------|
| 1 | 0-2058 | Malloy Lithographing Inc |
| 2 | 0-2060 | Malloy Lithographing Inc |
| 3 | 0-2436 | Malloy Lithographing Inc |
| 4 | 1-200-5164 | Federal Express Corporation |
| 5 | 1-202-2978 | Federal Express Corporation |

(114 rows selected)

Implicit inner join:

- No keyword JOIN in FROM clause
- MUST have join condition in WHERE clause

General Implicit syntax for inner join

```
SELECT select_list
FROM table_1, table_2 [, table_3]...
WHERE table_1.column_name operator table_2.column_name
      [AND table_2.column_name operator table_3.column_name]...
```

Implicit syntax that inner-joins four tables

```
SELECT vendor_name, invoice_number, invoice_date,  
       line_item_amt, account_description  
FROM   vendors v, invoices i, invoice_line_items li,  
       general_ledger_accounts gl  
WHERE  v.vendor_id = i.vendor_id  
       AND i.invoice_id = li.invoice_id  
       AND li.account_number = gl.account_number  
       AND (invoice_total - payment_total - credit_total) > 0  
ORDER BY vendor_name, line_item_amt DESC
```

The result set

| | VENDOR_NAME | INVOICE_NUMBER | INVOICE_DATE | LINE_ITEM_AMT | ACCOUNT_DESCRIPTION |
|---|-------------------------------|----------------|--------------|---------------|-----------------------|
| 1 | Abbey Office Furnishings | 203339-13 | 02-MAY-14 | 17.5 | Office Supplies |
| 2 | Blue Cross | 547481328 | 20-MAY-14 | 224 | Group Insurance |
| 3 | Blue Cross | 547480102 | 19-MAY-14 | 224 | Group Insurance |
| 4 | Blue Cross | 547479217 | 17-MAY-14 | 116 | Group Insurance |
| 5 | Cardinal Business Media, Inc. | 134116 | 01-JUN-14 | 90.36 | Card Deck Advertising |

(44 rows selected)

The implicit syntax for an outer join

```
SELECT select_list
FROM table_1, table_2 [, table 3]...
WHERE table_1.column_name [(+)] comp_operator
      table_2.column_name [(+)]
      [table_2.column_name [(+)] comp_operator
      table_3.column_name [(+)] ]...
```

- NO keyword JOIN in FROM clause
- Place (+) next to the table with missing (less) information in join condition
 - E.g. table1.column_name = table2.column_name (+)
 - + means: we need to ADD some information to table 2 for a matching value in table 1
- **CanNOT** place (+) on both sides of a comparison operator
 - So NO Implicit Full outer join in Oracle.

Implicit syntax with a left outer join

```
SELECT department_name AS dept_name,  
       dpt.department_number AS dept_no,  
       last_name  
FROM departments dpt, employees emp  
WHERE dpt.department_number = emp.department_number (+)  
ORDER BY department_name
```

The result set

| | DEPT_NAME | DEPT_NO | LAST_NAME |
|---|-------------|----------|-----------|
| 1 | Accounting | 1 | Hernandez |
| 2 | Maintenance | 5 | Hardy |
| 3 | Operations | 3 (null) | |
| 4 | Payroll | 2 | Simonian |
| 5 | Payroll | 2 | Aaronsen |
| 6 | Payroll | 2 | Smith |
| 7 | Personnel | 4 | Jones |
| 8 | Personnel | 4 | O'Leary |

Explicit LEFT OUTER JOIN Equivalent to Previous Slide

- Use keyword: LEFT JOIN

```
SELECT department_name, AS dept_name,  
       dpt.department_number AS dept_no,  
       last_name  
FROM   departments dpt LEFT JOIN employees emp  
       ON dpt.department_number =  
          emp.department_number  
ORDER BY department_name;
```

Implicit syntax with a right outer join

```
SELECT department_name AS dept_name,  
       emp.department_number AS dept_no,  
       last_name  
FROM departments dpt, employees emp  
WHERE dpt.department_number (+) = emp.department_number  
ORDER BY department_name
```

The result set

(9 rows selected)

| | DEPT_NAME | DEPT_NO | LAST_NAME |
|---|-------------|---------|-----------|
| 1 | Accounting | 1 | Hernandez |
| 2 | Maintenance | 5 | Hardy |
| 3 | Payroll | 2 | Aaronsen |
| 4 | Payroll | 2 | Simonian |
| 5 | Payroll | 2 | Smith |
| 6 | Personnel | 4 | Jones |
| 7 | Personnel | 4 | O'Leary |
| 8 | (null) | 6 | Locario |
| 9 | (null) | 6 | Watson |

Explicit RIGHT OUTER JOIN Equivalent to Previous Slide

- Use keyword: RIGHT JOIN

```
SELECT department_name, AS dept_name,  
       dpt.department_number AS dept_no,  
       last_name  
FROM   departments dpt RIGHT JOIN employees emp  
       ON dpt.department_number =  
          emp.department_number  
ORDER BY department_name;
```

A SELECT statement with JOIN and the USING keyword

```
SELECT invoice_number, vendor_name
FROM vendors
      JOIN invoices USING (vendor_id)
ORDER BY invoice_number
```

Join tables based on matching values in specified common (i.e. same name) column(s)

The result set

| | INVOICE_NUMBER | VENDOR_NAME |
|---|----------------|-----------------------------|
| 1 | 0-2058 | Malloy Lithographing Inc |
| 2 | 0-2060 | Malloy Lithographing Inc |
| 3 | 0-2436 | Malloy Lithographing Inc |
| 4 | 1-200-5164 | Federal Express Corporation |

(114 rows selected)

JOIN ON Equivalent to Previous Slide

```
SELECT invoice_number, vendor_name
FROM vendors v
      JOIN invoices i
      ON v.vendor_id = i.vendor_id
ORDER BY invoice_number;
```

The syntax for a JOIN with the USING keyword

```
SELECT select_list
FROM table_1
    [{LEFT|RIGHT|FULL} [OUTER]] JOIN table_2
        USING(join_column_1[, join_column_2]...)
    [[{LEFT|RIGHT|FULL} [OUTER]] JOIN table_3
        USING (join_column_2[, join_column_2]...)]...
```

- [A]: A is optional
- {A|B}: A or B

More Complex SELECT with JOIN and USING

```
SELECT department_name AS dept_name, last_name,  
project_number  
FROM departments  
    JOIN employees USING (department_number)  
    LEFT JOIN projects USING (employee_id)  
ORDER BY department_name
```

The result set

Note: for self-reading, not required in the course

| | DEPT_NAME | LAST_NAME | PROJECT_NUMBER |
|---|-------------|-----------|----------------|
| 1 | Accounting | Hernandez | P1011 |
| 2 | Maintenance | Hardy | (null) |
| 3 | Payroll | Simonian | P1012 |
| 4 | Payroll | Smith | P1012 |
| 5 | Payroll | Aaronsen | P1012 |
| 6 | Personnel | Jones | (null) |
| 7 | Personnel | O'Leary | P1011 |

(7 rows selected)

NATURAL JOIN

- Join two tables based on **matching values** in **ALL columns** in the two tables that have the **same names**.
 - **Don't need** to **explicitly specify** what columns to use in natural join
- **Only works correctly** if the database is **designed** in a **certain way**.
 - Natural join may not work if database design changes
- May have unexpected results for complex queries
- More common to use JOIN ON or JOIN USING

A SELECT statement with JOIN and the NATURAL keyword

```
SELECT invoice_number, vendor_name  
FROM vendors  
      NATURAL JOIN invoices  
ORDER BY invoice_number
```

The result set

| | INVOICE_NUMBER | VENDOR_NAME |
|---|----------------|-----------------------------|
| 1 | 0-2058 | Malloy Lithographing Inc |
| 2 | 0-2060 | Malloy Lithographing Inc |
| 3 | 0-2436 | Malloy Lithographing Inc |
| 4 | 1-200-5164 | Federal Express Corporation |

(114 rows selected)

Natural JOIN - Join two tables based on matching values in ALL columns in the two tables that have the same names

JOIN ON Equivalent to Previous Slide

- Only 1 common column vendor_id

```
SELECT invoice_number, vendor_name
FROM vendors v
      JOIN invoices i
      ON v.vendor_id = i.vendor_id
ORDER BY invoice_number;
```

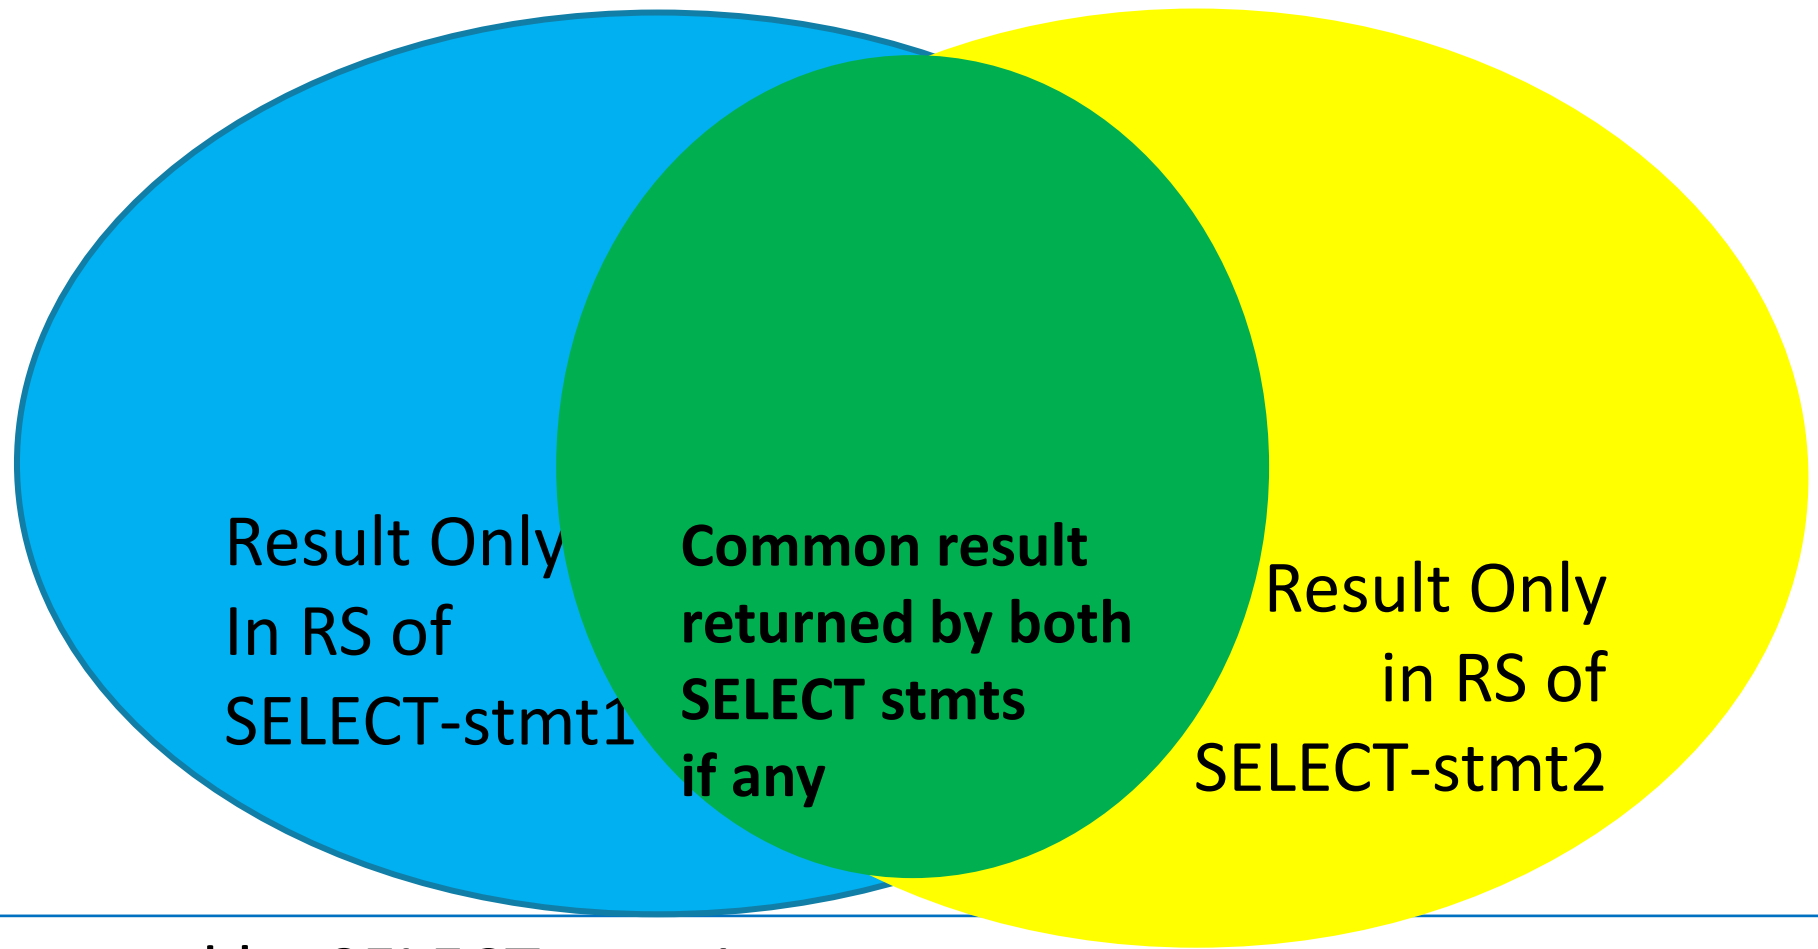

Set Operators in SQL

- Used to combine query results from multiple SELECT statements
- Each SELECT statement returns a set of rows: ResultSet.
- Set operators in SQL
 - UNION, INTERSECT, MINUS
- UNION operator
 - Same as set operator \cup in Discrete Math
- INTERSECT operator
 - Same as set operator \cap in Discrete Math
- MINUS operator
 - Same as set operator $-$ in Discrete Math

Review: Set Operators in Discrete Math

- Sets $A = \{1, 2, 3\}$, $B = \{3, 4, 5\}$
- $A \cup B = \{1, 2, 3, 4, 5\}$
- $A \cap B = \{3\}$
- $A - B = \{1, 2\}$
- $B - A = \{4, 5\}$

RS: result set
stmt: statement



Blue: result only returned by SELECT-stmt1

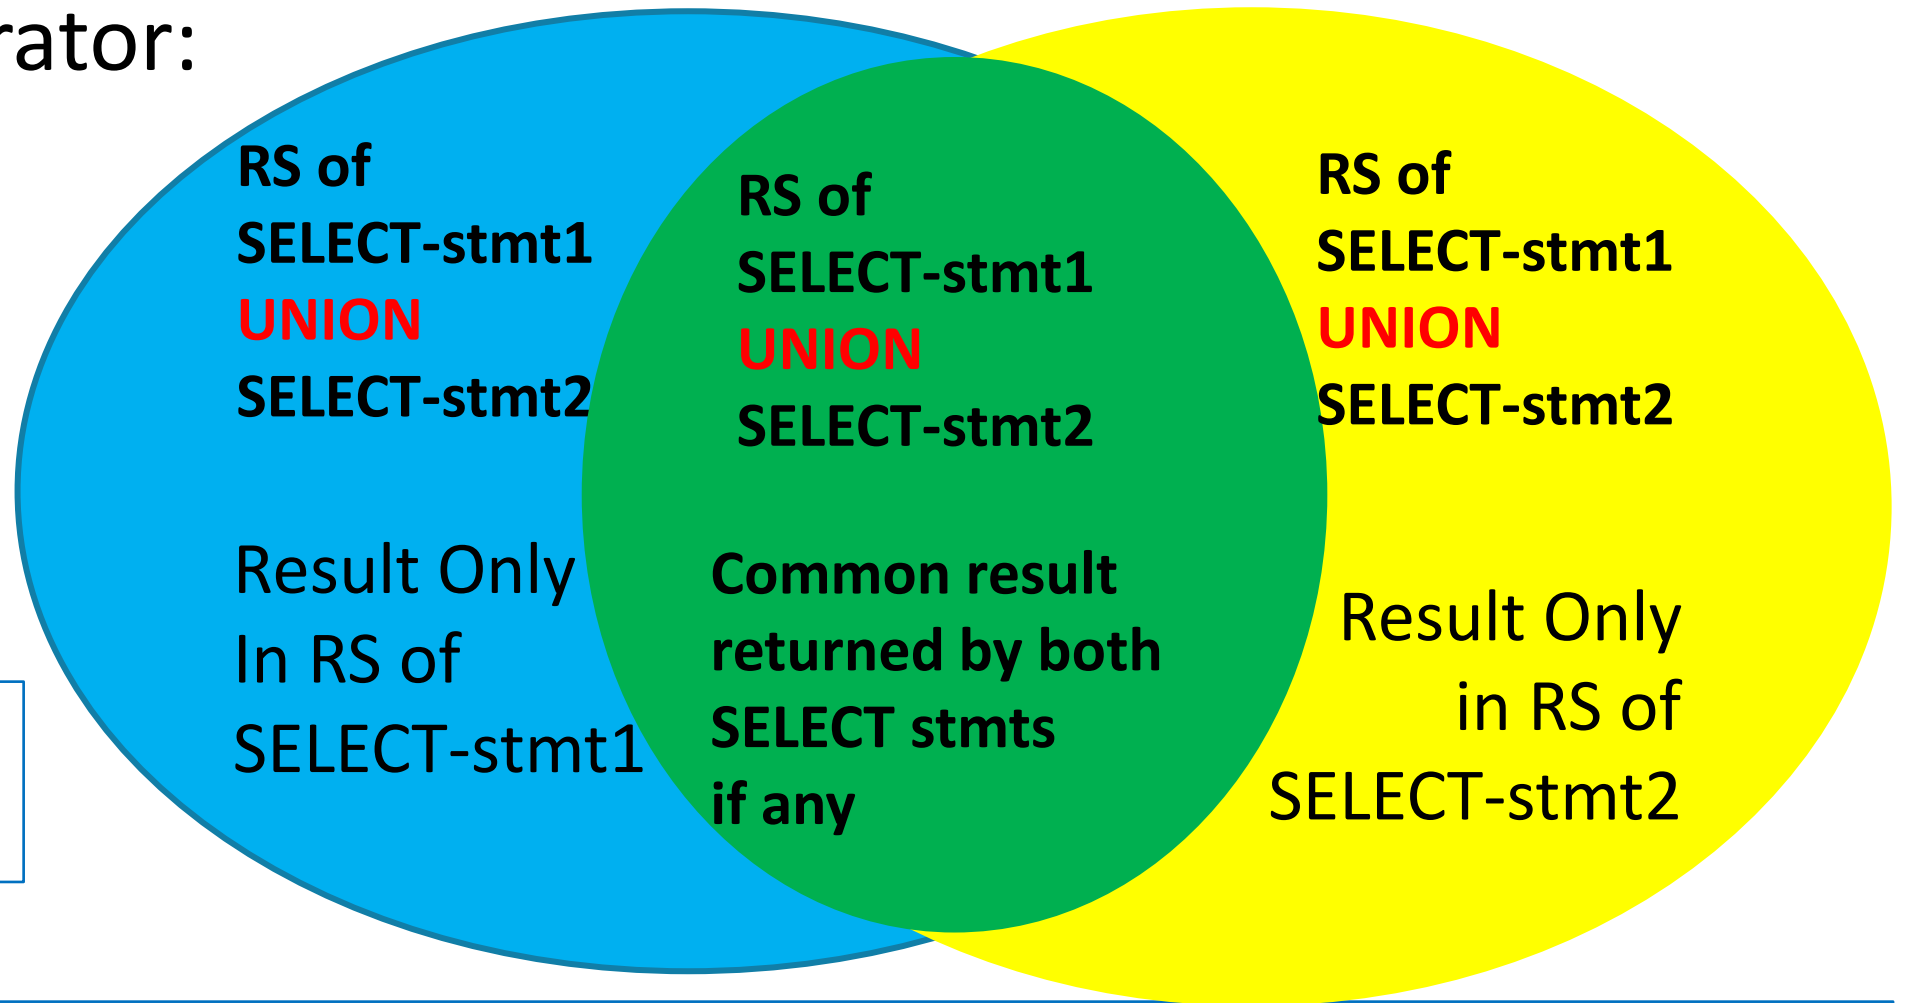
Yellow: result only returned by SELECT-stmt2

Green: result returned by both SELECT-stmt1 and SELECT-stmt2 if there is any.

Result Set of SELECT-stmt1 = Blue-Area + Green-Area

Result Set of SELECT-stmt2 = Yellow-Area + Green-Area

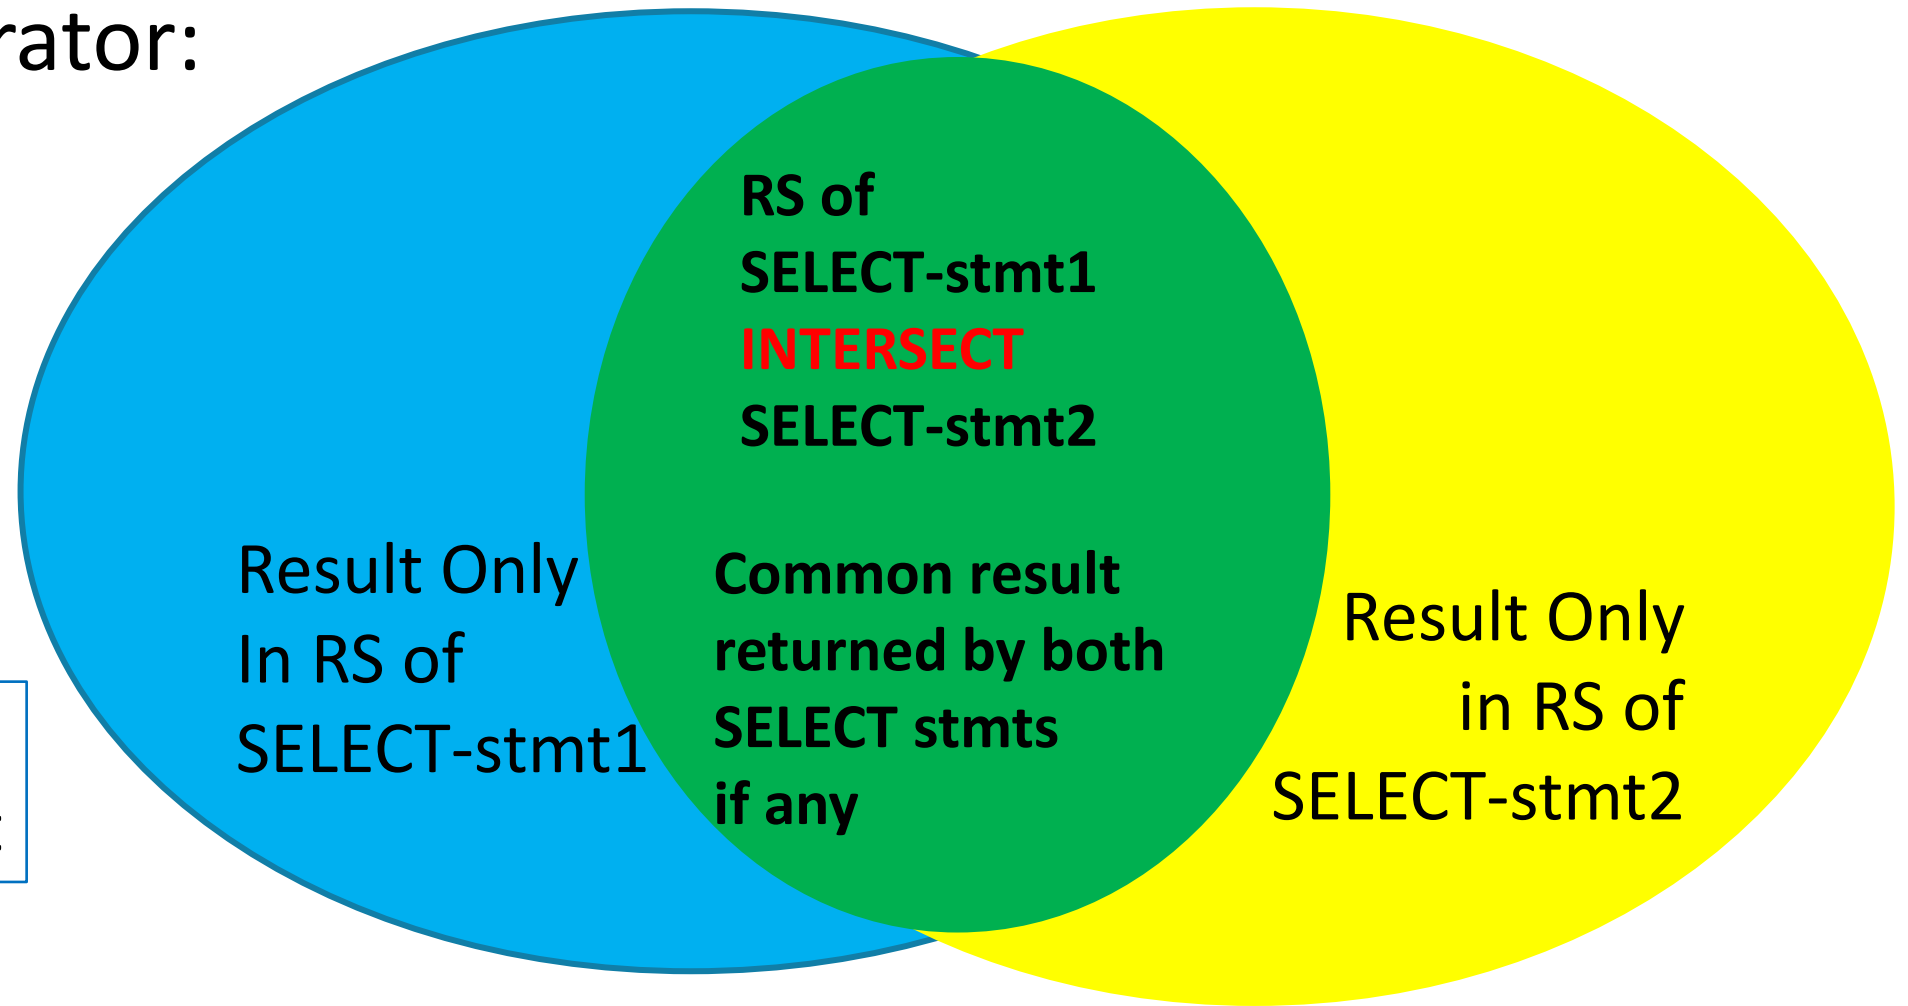
SQL set operator: UNION



RS: result set
stmt: statement

Result Set of (SELECT-stmt1 **UNION** SELECT-stmt2) = Blue + Green + Yellow
= ALL in Result Set of SELECT-stmt1 (Blue + Green)
+ ALL in Result Set of SELECT-stmt2 (Yellow + Green) – **DuplicatesInGreen**

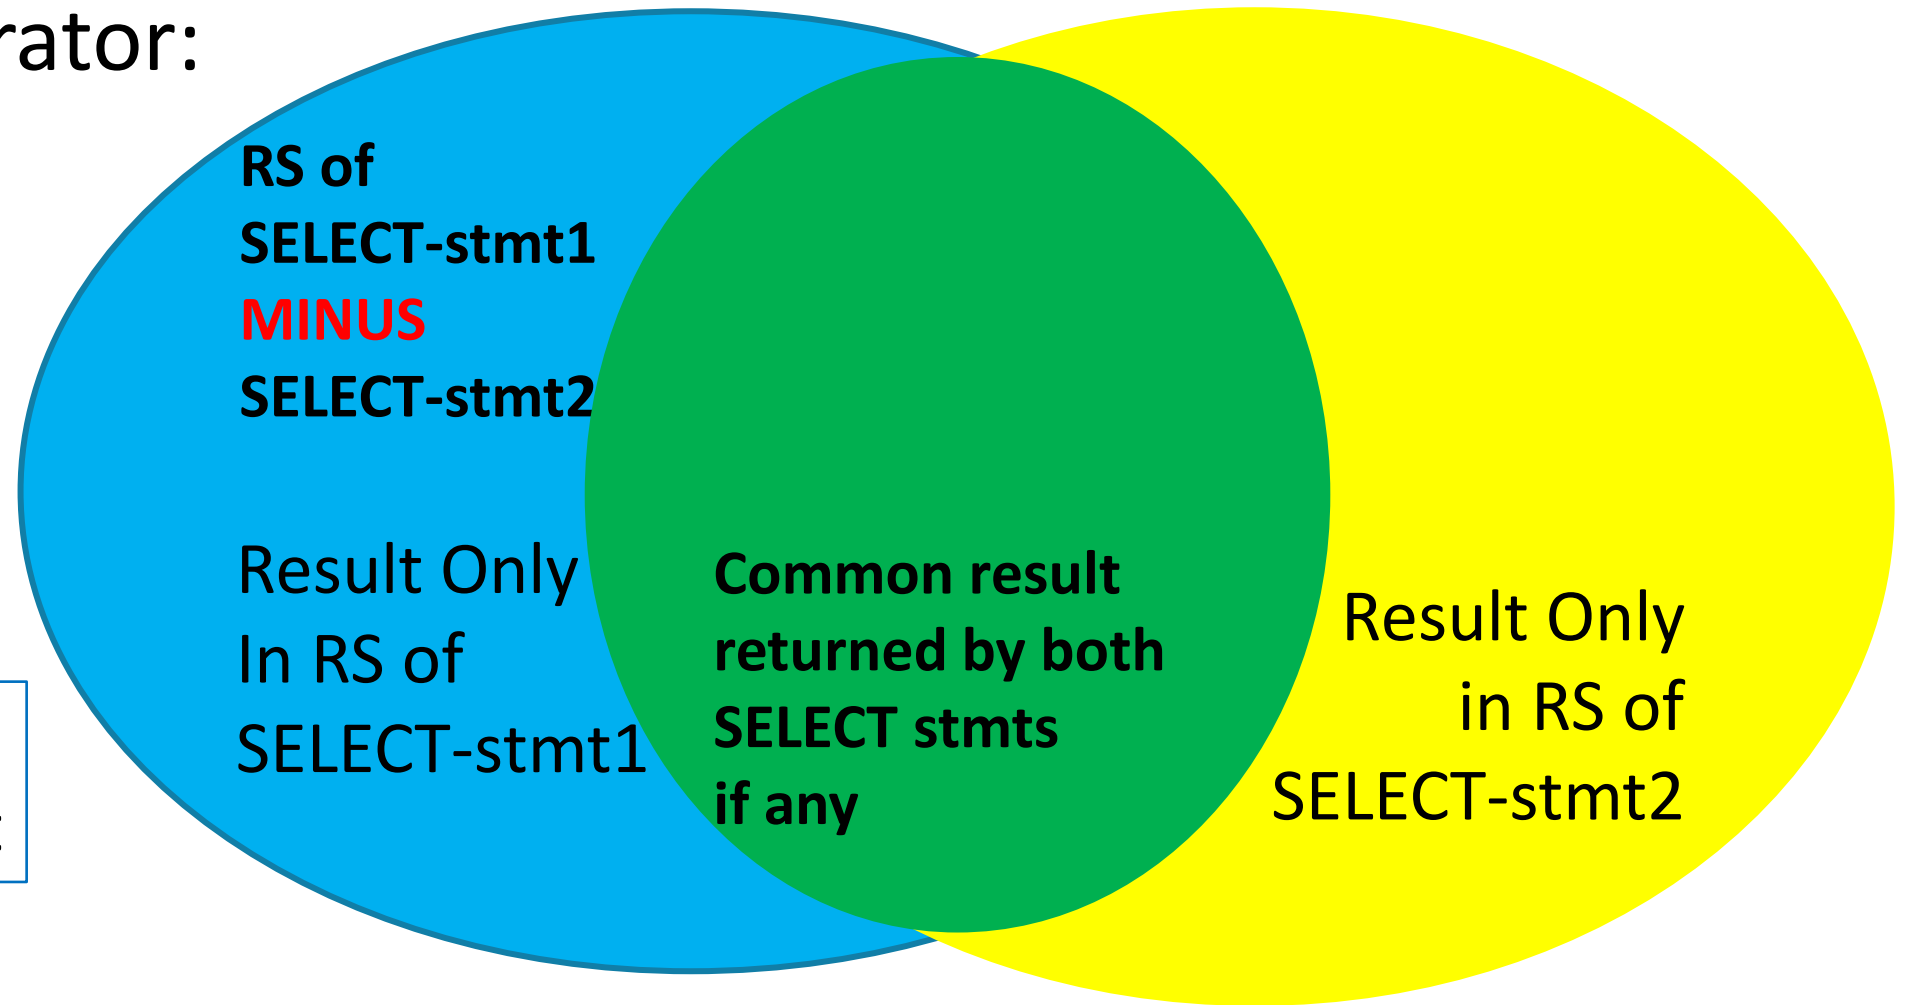
SQL set operator: INTERSECT



RS: result set
stmt: statement

Result Set of (SELECT-stmt1 **INTERSECT** SELECT-stmt2) = Green
= Common result in BOTH Result Set of SELECT-stmt1 (Blue + Green)
AND Result Set of SELECT-stmt2 (Yellow + Green)

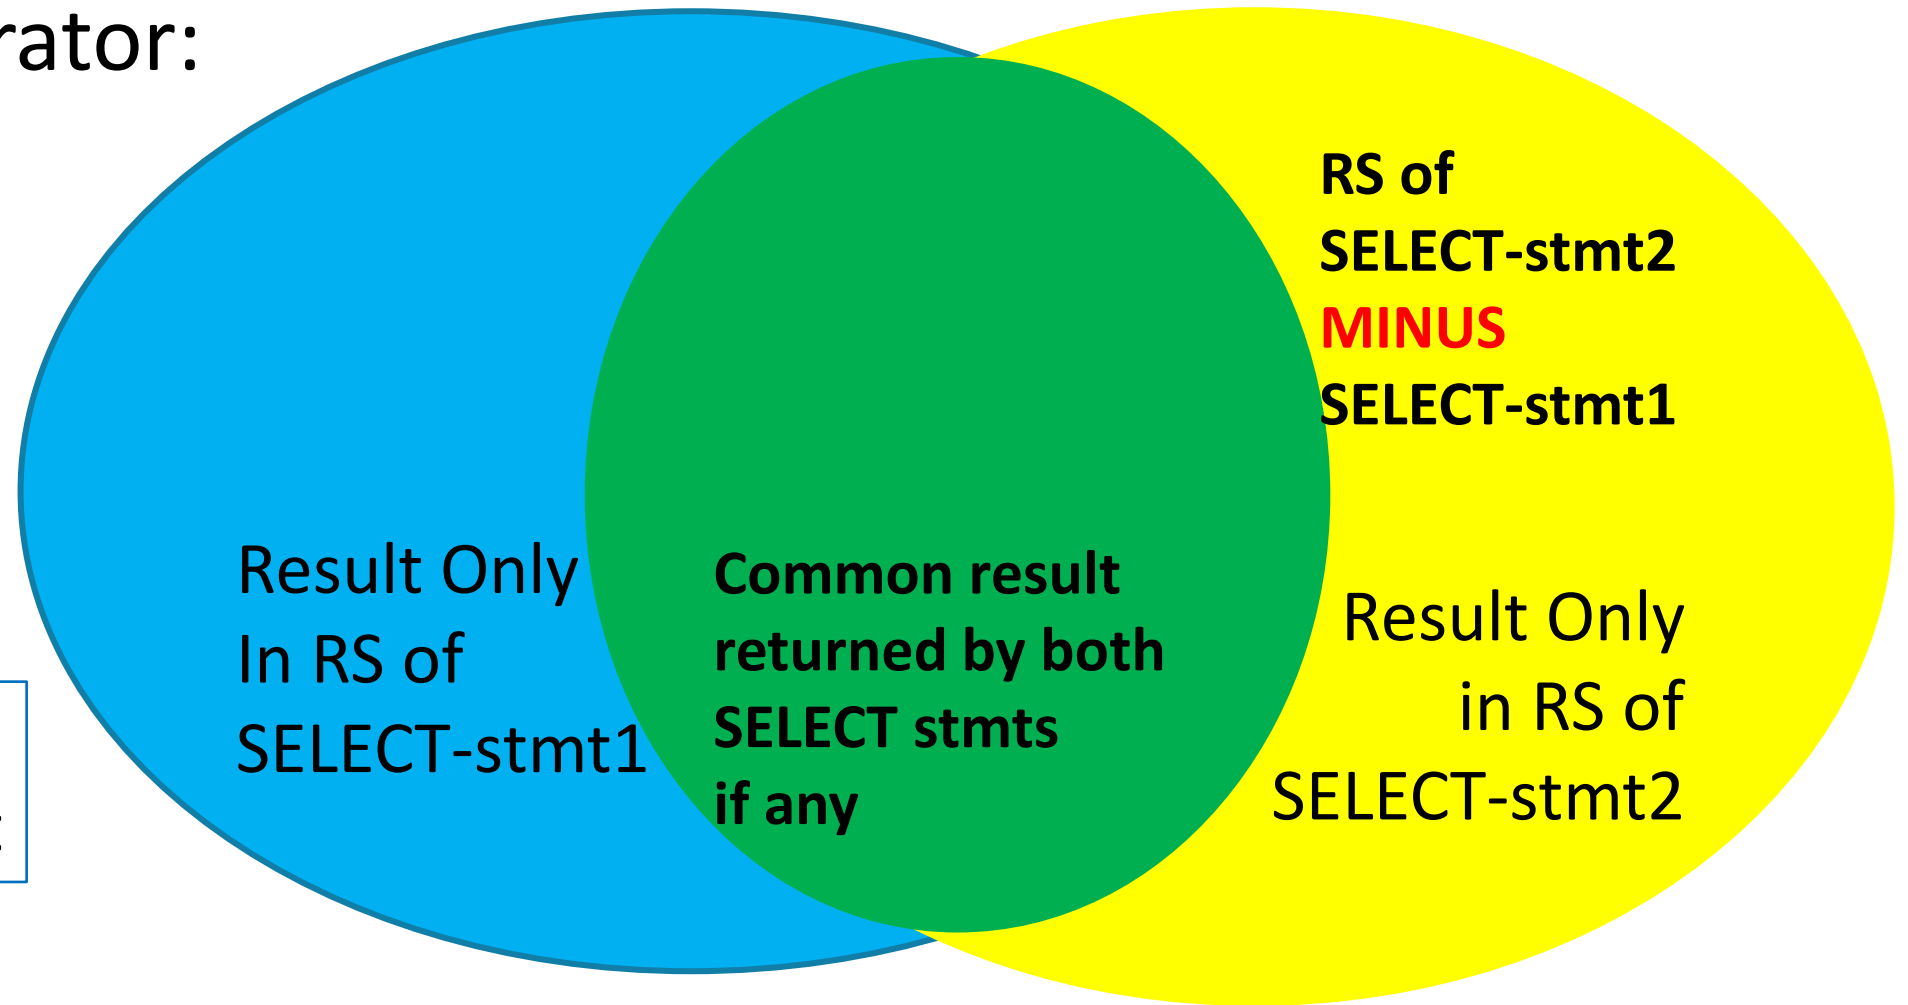
SQL set operator: MINUS



RS: result set
stmt: statement

Result Set of (SELECT-stmt1 **MINUS** SELECT-stmt2) = Blue
= Result ONLY in Result Set of SELECT-stmt1 (Blue + Green), but
NOT in Result Set of SELECT-stmt2 (Yellow + Green)

SQL set operator: MINUS



RS: result set
stmt: statement

Result Set of (SELECT-stmt2 **MINUS** SELECT-stmt1) = YELLOW
= Result ONLY in Result Set of SELECT-stmt2 (Yellow + Green), but
NOT in Result Set of SELECT-stmt1 (Blue + Green)

The Customers table

| ⚡ CUSTOMER_LAST_NAME | ⚡ CUSTOMER_FIRST_NAME |
|----------------------|-----------------------|
| Anders | Maria |
| Trujillo | Ana |
| Moreno | Antonio |
| Hardy | Thomas |
| Berglund | Christina |
| Moos | Hanna |

(24 rows selected)

The Employees table

| ⚡ EMPLOYEE_ID | ⚡ LAST_NAME | ⚡ FIRST_NAME | ⚡ DEPARTMENT_NUMBER |
|---------------|-------------|--------------|---------------------|
| 1 | Smith | Cindy | 2 |
| 2 | Jones | Elmer | 4 |
| 3 | Simonian | Ralph | 2 |
| 4 | Hernandez | Olivia | 1 |
| 5 | Aaronsen | Robert | 2 |
| 6 | Watson | Denise | 6 |
| 7 | Hardy | Thomas | 5 |
| 8 | O'Leary | Rhea | 4 |
| 9 | Locario | Paulo | 6 |

(9 rows selected)

A query that includes rows that occur in first query or second query

```
SELECT customer_first_name, customer_last_name
FROM customers
UNION
SELECT first_name, last_name
FROM employees
ORDER BY customer_last_name
```

Different column names in result
sets of two SELECT statements

The result set

| | CUSTOMER_FIRST_NAME | CUSTOMER_LAST_NAME |
|---|---------------------|--------------------|
| 1 | Robert | Aaronsen |
| 2 | Maria | Anders |
| 3 | Christina | Berglund |
| 4 | Art | Braunschweiger |
| 5 | Donna | Chelan |
| 6 | Fred | Citeaux |
| 7 | Thomas | Hardy |
| 8 | Olivia | Hernandez |

(32 rows selected)

A query that **only includes rows** that **occur in both queries**

```
SELECT customer_first_name, customer_last_name
FROM customers
INTERSECT
SELECT first_name, last_name
FROM employees
ORDER BY customer_last_name
```

Different column names in
result sets of two SELECT
statements

The result set

| | ⚡ CUSTOMER_FIRST_NAME | ⚡ CUSTOMER_LAST_NAME | |
|---|-----------------------|----------------------|--|
| 1 | Thomas | Hardy | |

(1 rows selected)

A query that **excludes rows from the first query** **if they also occur in the second query**

```
SELECT customer_first_name, customer_last_name  
FROM customers
```

MINUS

```
SELECT first_name, last_name  
FROM employees  
ORDER BY customer_last_name
```

Different column names in result
sets of two SELECT statements

The result set

| | ⚡ CUSTOMER_FIRST_NAME | ⚡ CUSTOMER_LAST_NAME |
|---|-----------------------|----------------------|
| 1 | Maria | Anders |
| 2 | Christina | Berglund |
| 3 | Art | Braunschweiger |
| 4 | Donna | Chelan |

(23 rows selected)

The syntax for a union

```
SELECT_statement_1
UNION [ALL]
SELECT_statement_2
[UNION [ALL]
SELECT_statement_3] ...
[ORDER BY order_by_list]
```

[A]: A is optional

ALL - Include duplicates in result set of UNION
Default: no duplicate

Rules for a union

- The **number of columns** must be the **same** in all SELECTs.
- The **column data types** must be **compatible**.
- The **column names** in the result are taken **from the first SELECT** statement.
- Column names in **ORDER BY** must be **same as** those **in first SELECT** statement.

Different SELECT statements may have different column names in their SELECT clauses.

The syntax for **MINUS** and **INTERSECT** operations

```
SELECT_statement_1  
{MINUS | INTERSECT}  
SELECT_statement_2  
[ORDER BY order_by_list]
```

- [A]: A is optional
- {A | B}: A or B

Same rules as for a union

- The number of columns must be the same in all SELECTs.
- The column data types must be compatible.
- The column names in the result are taken from the first SELECT statement.
- Column names in ORDER BY must be same as those in first SELECT statement. (Can also use column position in ORDER BY)

Different SELECT statements may have different column names in their SELECT clauses.